

Homework 7, MATH 455: Due Mon, 04/30/2018

Alexander Van Roijen

May 1, 2018

Instructions: The homework assignment editing this L^AT_EX document. Download the L^AT_EX source from the class web page and study it to learn more about L^AT_EX. Replace the text with appropriate information. Run “pdflatex” on this document.

You will submit this assignment in two parts:

1. Print out the PDF file and bring it to class, and
2. Send an e-mail to:

gang@math.binghamton.edu

before class on the due date with two attachments:

- The L^AT_EX source file, and
- The generated PDF document.

Please complete the following:

1. Finish R exercises 11.1, 11.2, 11.3, 11.4, 11.6 of the textbook. Submit your answers for **ALL** questions.

- (a) 11.1 We first take a look at the PC

```
> hold=prcomp(seatpos[,-c(9,1,2)])  
> print(summary(hold))
```

Importance of components:

PC1	PC2	PC3	PC4	PC5	PC6	
Standard deviation	17.1573	2.89689	2.11907	1.56412	1.22502	0.46218
Proportion of Variance	0.9453	0.02695	0.01442	0.00786	0.00482	0.00069
Cumulative Proportion	0.9453	0.97222	0.98664	0.99450	0.99931	1.00000

Looks like the first two components explain most of the variation of our data. Using them for our prediction we get the following.

```
> cmonnow = pcr(hipcenter~.-Age-Weight,data=seatpos[,ncomp=2])  
> predict(cmonnow,testhcf,ncomp=2,interval="prediction")  
  
, , 2 comps
```

```
hipcenter  
1 -204.4636
```

- (b) 11.2 We fit a partial least squares model to the same data and examine the number of components to use.




Figure 1: examining the residual mean squared error over number of components and choosing the min value

we then use 4 components as it has the minimum RMSEP value and we get the following prediction

```
> splsmod <- plsr(hipcenter ~ ., data=seatpos, validation="CV")
> #4 components looks good
> hcpred = predict(splsmod, testhcf, ncomp=4)
> print(hcpred)
, , 4 comps
```

```
hipcenter
1 -179.4634
```

(c) 11.3 We are now going to fit a ridge regression model to the seatpos data




Figure 2: examining th

using the minimum lambda value provided of 0.05, we get the following prediction

```
> hcrgpred1 = cbind(1,as.matrix(testhcf[1,]))%*%coef(hcrgmod2)[8,]
> hcrgpred1
[,1]
1 -175.488
```

(d) 11.4

We first remove each tenth observation and separate the data.

```
fat2=fat[-seq(1,length(fat[,1]),10),]
testfat = fat[seq(1,length(fat[,1]),10),]
```

- i. a we now fit a linear model and get the following prediction accuracy described by the residual mean squared error between the predictions and the actual observations

```
> oglg = lm(siri ~ . -brozek -density,fat2)
```

```
> wut=predict(oglg,newdata=testfat)
> rmse(wut,testfat$siri)
[1] 1.946023
```

ii. b we now use the stepwise function to determine the "ideal" model

```
> stepwise(lm(siri ~ . -brozek -density,fat2),criterion = c("AIC"),direction=c("
Call:
```

```
lm(formula = siri ~ abdom + free + weight + forearm + adipos +
thigh + chest + biceps + ankle, data = fat2)
```

Coefficients:

(Intercept)	abdom	free	weight	forearm	adipos	ankle
-2.9190	0.1179	-0.5698	0.3925	0.2146	-0.5277	0.1475

I chose forward progression, and proceeded to fit a model with the chosen parameters and got the following prediction results

```
> splg = lm(formula = siri ~ abdom + free + weight + forearm + adipos + thigh +
> wut2=predict(splg,newdata=testfat)
> rmse(wut2,testfat$siri)
[1] 1.98911
```

we see a slightly higher RMSE, but overall quite close and simpler too

iii. c Now we want to fit a principle component regression onto our data.

```
> print(summary(temp))
```

Importance of components:

PC1	PC2	PC3	PC4	PC5	PC6	PC7	PC8	PC9	PC10
Standard deviation			36.8986	15.5341	10.29573	3.66009	3.44451	2.64961	2.14660
Proportion of Variance			0.7736	0.1371	0.06023	0.00761	0.00674	0.00399	0.00262
Cumulative Proportion			0.7736	0.9107	0.97095	0.97856	0.98531	0.98929	0.99191
PC13	PC14	PC15	PC16						
Standard deviation				1.06850	1.00511	0.75913	0.46948		

```
Proportion of Variance 0.00065 0.00057 0.00033 0.00013
Cumulative Proportion 0.99897 0.99955 0.99987 1.00000
```

I choose to only include the first 3 PRC as they cover about 97 percent of the variation in the data.

Fitting the model, we now get

```
> fatpcr = pcr(siri ~ . -brozek -density, data=fat2, ncomp=3)
> pcrr= predict(fatpcr, testfat, ncomp=3, interval="prediction")
> rmse(pcrr, testfat$siri)
[1] 2.487871
```

the RMSE is quite higher, looking at a PCR with all PCs we get

```
> rmse(pcrr2, testfat$siri)
[1] 1.946023
```

so we could improve our RMSE, but that would effectively negate the point of PCR.

iv. d Looking at a partial least squares regression, we create the following




Figure 3: a look at at the ideal number of components for our PLSR

We determine graphically and within the code that 4 components is the ideal choice. We achieve the following.

```
> summary(fatpls)
```

```
Data: X dimension: 226 15
```

```
Y dimension: 226 1
```

```
Fit method: kernelpls
```

```
Number of components considered: 15
```

```
VALIDATION: RMSEP
```

```
Cross-validated using 10 random segments.
```

(Intercept)	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps	7 comps	8 co
CV	8.382	6.384	2.998	1.967	1.709	1.822	1.889	1.95
adjCV	8.382	6.379	2.997	1.962	1.698	1.800	1.858	1.92
12 comps	13 comps	14 comps	15 comps					

CV	2.848	2.906	2.944	2.952
adjCV	2.749	2.804	2.839	2.846

TRAINING: % variance explained

	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps	7 comps	8 comps	9 comps
X	76.85	90.81	97.13	97.83	98.24	98.54	99.04	99.27	
siri	44.30	87.73	95.02	96.72	96.94	97.07	97.11	97.14	

	13 comps	14 comps	15 comps
X	99.83	99.92	100.00
siri	97.16	97.16	97.16

```
> hcpred = predict(fatpls,testfat,ncomp=4)
```

```
> rmse(hcpred,testfat$siri)
```

```
[1] 1.973459
```

Similar predictive ability to other models we have seen

- v. e Ridge regression! Exciting stuff, this bad boy will let us trim down the size of our model a bit by utilizing a penalty term! Lets examine the potentials of this idea.




Figure 4: a look at at the ideal number of components for our PLSR

This is hard to interpret, so lets analyze the minimum GCV to get an idea of what our penalty term should be.

```
> fatrm = lm.ridge(siri ~ .-brozek -density,data=fat2,lambda=seq(0,2,len=30))
> which.min(fatrm$GCV)
0.00000000
1
```

This zero implies that our best bet is zero penalty, which equivalently that reducing the model provides no optimization here and the OLS is best. We find that the RMSE is the same as our OLS in part a, which coincides with our thoughts.

```
> hcrpred1 = cbind(1,as.matrix(rmtestfat))%*(coef(fatrm)[1,])
> rmse(hcrpred1,testfat$siri)
[1] 1.946023
```

(e) 11.6

i. a

```

> withoutData=kanga[complete.cases(kanga),]
> withoutData$sex = as.numeric(as.factor(withoutData$sex))
> withoutData$species = as.numeric(as.factor(withoutData$species))
> temp = prcomp(withoutData)
> print(summary(temp))
Importance of components:
Importance of components:
PC1      PC2      PC3      PC4      PC5      PC6      PC7
Standard deviation      288.0383 69.51358 30.74721 27.85652 21.73040 19.42447 17.
Proportion of Variance  0.9002  0.05243  0.01026  0.00842  0.00512  0.00409  0.
Cumulative Proportion  0.9002  0.95268  0.96294  0.97136  0.97648  0.98058  0.

```

As we can see, our first principal component covers 90 percent of variation

ii. b

```

> print(temp$rotation[,1])
species                sex      basilar.length occipitonasal.length      pa
-9.470452e-05      -6.133600e-04      -4.840682e-01      -4.562961e-01
palate.width      nasal.length      nasal.width      squamosal.depth
-8.435755e-02      -2.480989e-01      -7.464748e-02      -6.366132e-02
zygomatic.width      orbital.width      .rostral.width      occipital.depth
-2.066976e-01      -1.428888e-02      -1.064589e-01      -1.781770e-01
foramina.length      mandible.length      mandible.width      mandible.depth
-9.941184e-03      -4.359818e-01      -2.999679e-02      -5.832212e-02

```

As we can see, mandible length, occipitonasal length, palate length, and basilar length are the prominent terms

iii. c

```

> temp2 = prcomp(withoutData,scale=TRUE)
> print(summary(temp2))
Importance of components:
PC1      PC2      PC3      PC4      PC5      PC6      PC7      PC8      ...
Standard deviation      3.5509 1.5178 1.11386 1.00232 0.84943 0.69411 0.55590 0.5

```

Proportion of Variance	0.6304	0.1152	0.06203	0.05023	0.03608	0.02409	0.01545	0.0
Cumulative Proportion	0.6304	0.7456	0.80765	0.85788	0.89396	0.91805	0.93350	0.9

We can tell that, similar to the fat data set we analyzed in class, there is a leveled out effect. Large values are made less significant and thus the proportion of variance per component is less prominent.

iv. d

```
> print(temp2$rotation[,1])
```

species	sex	basilar.length	occipitonasal.length	pa
palate.width	nasal.length	nasal.width	squamosal.depth	
0.23270782	0.22767199	0.22577402	0.23719817	
zygomatic.width	orbital.width	.rostral.width	occipital.depth	
0.25677431	0.07539545	0.25588530	0.26958527	
foramina.length	mandible.length	mandible.width	mandible.depth	
0.05683491	0.27711781	0.20971711	0.23837398	

```
> print(temp2$rotation[,2])
```

species	sex	basilar.length	occipitonasal.length	pa
0.5649159512	-0.0459169262	0.0090950112	0.1659461872	
palate.width	nasal.length	nasal.width	squamosal.depth	
0.0121999313	0.3258713200	0.2396642898	-0.1554982206	
zygomatic.width	orbital.width	.rostral.width	occipital.depth	
-0.2214711336	0.0502588495	0.0007699419	-0.0243930064	
foramina.length	mandible.length	mandible.width	mandible.depth	
0.3370527404	-0.0189223029	-0.3231510620	-0.2003645746	

the second PC shows a strong correlation with species, crest width, nasal length, and formalin length. The first one represented that a lot of these kangaroos are a lot more similar, while the second PC, orthogonal to the first, shows that they differ in the corresponding high valued components. We can see this mainly some higher values, like .33 for foramina length in the second PC. Unlike the first, it is not as level, and thus different components **differentiate** the kangaroos

v. e




Figure 5: A mahalanobis distance graph for the kanga data set

We can see some pretty extreme values. It may be worth looking into these data points. They are likely very informational

- vi. f I was not sure what to do with this, I have a graph of the two principal components plotted against each other, but I am not sure how you would differentiate the components by sex?




Figure 6: A plot of the coefficients of PC1 against PC2

I suppose we can see some orthogonality / differences between components , as we would expect.

2. Finish R exercises 13.2, 13.3 of the textbook. Submit your answers for **ALL** questions.

(a) 13.2

i. a

```
> glm = lm(Species~.-Endemics,gala)
> summary(glm)
```

Call:

```
lm(formula = Species ~ . - Endemics, data = gala)
```

Residuals:

Min	1Q	Median	3Q	Max
-111.679	-34.898	-7.862	33.460	182.584

Coefficients:

```
Estimate Std. Error t value Pr(>|t|)
(Intercept)  7.068221  19.154198   0.369 0.715351
Area          -0.023938   0.022422  -1.068 0.296318
Elevation      0.319465   0.053663   5.953 3.82e-06 ***
Nearest        0.009144   1.054136   0.009 0.993151
Scruz          -0.240524   0.215402  -1.117 0.275208
Adjacent       -0.074805   0.017700  -4.226 0.000297 ***
---
Signif. codes:  0 *** 0.001 ** 0.01 * 0.05 . 0.1 1
```

Residual standard error: 60.98 on 24 degrees of freedom

Multiple R-squared: 0.7658, Adjusted R-squared: 0.7171

F-statistic: 15.7 on 5 and 24 DF, p-value: 6.838e-07

Here is our simple model with a summary with the 5 geographic predictors

ii. b

```
> gmlm = lm(NS~Area+Anear+Dist+DistSC+Elevation,galamiss)
> summary(gmlm)
```

Call:

```
lm(formula = NS ~ Area + Anear + Dist + DistSC + Elevation, data = galamiss)
```

Residuals:

```
Min      1Q  Median      3Q      Max
-114.13 -38.90 -10.03   35.34  172.19
```

Coefficients:

```
Estimate Std. Error t value Pr(>|t|)
(Intercept) 17.99073  29.89638   0.602 0.555269
```

Area	-0.02700	0.02637	-1.024	0.320243
Anear	-0.07822	0.02159	-3.623	0.002103 **
Dist	-0.09376	1.21083	-0.077	0.939182
DistSC	-0.29841	0.26619	-1.121	0.277857
Elevation	0.32213	0.06766	4.761	0.000181 ***

Signif. codes: 0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

Residual standard error: 69.17 on 17 degrees of freedom

(6 observations deleted due to missingness)

Multiple R-squared: 0.7636, Adjusted R-squared: 0.6941

F-statistic: 10.98 on 5 and 17 DF, p-value: 7.594e-05

As we can see, missing values have reduce the number of observations available, and lowered our degrees of freedom. Our confidence has thus decreased compared to the OG.

iii. c

```
> gmmeans = colMeans(galamiss,na.rm = TRUE)
> imgalamiss = galamiss
> for(i in c(2:8)) imgalamiss[is.na(galamiss[,i]),i] <- gmmeans[i]
> gmimlm = lm(NS~Area+Anear+Dist+DistSC+Elevation,imgalamiss)
> summary(gmimlm)
```

Call:

```
lm(formula = NS ~ Area + Anear + Dist + DistSC + Elevation, data = imgalamiss)
```

Residuals:

Min	1Q	Median	3Q	Max
-96.00	-45.43	-11.11	28.64	223.83

Coefficients:

```

Estimate Std. Error t value Pr(>|t|)
(Intercept) -13.076462 31.291432 -0.418 0.67990
Area          0.000602  0.027810  0.022 0.98292
Anear        -0.064403  0.023002 -2.800 0.01017 *
Dist          0.403334  1.327801  0.304 0.76404
DistSC       -0.077887  0.285100 -0.273 0.78714
Elevation     0.269094  0.072546  3.709 0.00115 **
---
Signif. codes:  0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

```

```

Residual standard error: 77.03 on 23 degrees of freedom
Multiple R-squared:  0.6382, Adjusted R-squared:  0.5595
F-statistic: 8.113 on 5 and 23 DF,  p-value: 0.0001563

```

our R squared has gone down significantly , I expect it is due to the imputed mean values to mean very little to our model and actually skew it, like we have seen in class. I looked at the correlation just to see if there was any to worry about.

```

> cor(imgalamiss)

```

NS	ES	Area	Anear	Dist	DistSC	Elevation
NS	1.00000000	0.973531485	0.6162212	0.02055711	-0.019642065	-0.186226
ES	0.97353148	1.000000000	0.6177911	0.07370982	-0.003761533	-0.181223
Area	0.61622115	0.617791136	1.0000000	0.17735771	-0.114527566	-0.108972
Anear	0.02055711	0.073709818	0.1773577	1.0000000	-0.119674948	0.044988
Dist	-0.01964207	-0.003761533	-0.1145276	-0.11967495	1.000000000	0.614269
DistSC	-0.18622612	-0.181223511	-0.1089720	0.04498817	0.614269601	1.000000
Elevation	0.67837148	0.723529575	0.7433463	0.54805436	-0.105879903	-0.158133
EM	0.28517661	0.310320605	0.1575627	0.06360495	0.293824516	0.416229

```

> cor(galamiss)

```

NS	ES	Area	Anear	Dist	DistSC	Elevation
NS	1.00000000	0.973531485	0.6162212	0.02055711	-0.019642065	-0.186226

ES	0.97353148	1.000000000	0.6177911	0.07370982	-0.003761533	-0.181223
Area	0.61622115	0.617791136	1.0000000	0.17735771	-0.114527566	-0.108972
Anear	0.02055711	0.073709818	0.1773577	1.00000000	-0.119674948	0.044988
Dist	-0.01964207	-0.003761533	-0.1145276	-0.11967495	1.000000000	0.614269
DistSC	-0.18622612	-0.181223511	-0.1089720	0.04498817	0.614269601	1.000000
Elevation	NA	NA	NA	NA	NA	NA
EM	0.28517661	0.310320605	0.1575627	0.06360495	0.293824516	0.416229

We can see that elevation is the missing data, which was our most valuable predictor in our original models. by adding mean imputed values, it has diminished the predictors ability to explain the variance in the model. There is also some correlation added by providing complete data for the model, but its not ridiculous... yet...

iv. d

```

> elevlm = lm(Elevation~Area+Anear+Dist+DistSC,galamiss)
> rggalamiss=galamiss
> galamiss[is.na(galamiss$Elevation),]
NS ES  Area  Anear  Dist  DistSC  Elevation  EM
Baltra      58 23 25.09   1.84  0.6    0.6      NA  0
Coamano      2  1  0.05 903.82  1.9    1.9      NA  0
Daphne_Major 18 11  0.34   1.84  8.0    8.0      NA  0
Eden         8  4  0.03  17.95  0.4    0.4      NA  0
Las_Plazas   12  9  0.23  25.09  0.5    0.6      NA  0
Seymour     44 16  1.84  25.09  0.6    9.6      NA  0
> rgvals = predict(elevlm,galamiss[is.na(galamiss$Elevation),])
> sepcounter=1
> for(i in 1:length(galamiss[,1])){
+   if(is.na(rggalamiss[i,7]))
+   {
+     rggalamiss[i,7]=rgvals[sepcounter]
+     sepcounter = sepcounter+1

```

```

+   }
+ }

> rgimlm = lm(NS~Area+Anear+Dist+DistSC+Elevation,rggalamiss)
> summary(rgimlm)

Call:
lm(formula = NS ~ Area + Anear + Dist + DistSC + Elevation, data = rggalamiss)

Residuals:
    Min       1Q   Median       3Q      Max
-110.26  -30.88  -14.52   29.51  197.48

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) -14.83873    26.26329  -0.565  0.57754
Area         -0.01876     0.02610  -0.719  0.47942
Anear        -0.07829     0.02137  -3.664  0.00129 **
Dist          0.05585     1.19947   0.047  0.96326
DistSC       -0.09669     0.25188  -0.384  0.70459
Elevation     0.32213     0.06757   4.767 8.31e-05 ***
---
Signif. codes:  0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

Residual standard error: 69.07 on 23 degrees of freedom
Multiple R-squared:  0.7091, Adjusted R-squared:  0.6459
F-statistic: 11.21 on 5 and 23 DF,  p-value: 1.459e-05

```

I wanted to see which values were missing information first. I then applied the regressed values to replace the NA points. We get a boosted performance! However, I suspect that our correlation will have increased compared to both the

missing data and the mean imputed values.

```
> cor(rggalamiss)
```

NS	ES	Area	Anear	Dist	DistSC	Elevation
NS	1.00000000	0.973531485	0.6162212	0.02055711	-0.019642065	-0.186226
ES	0.97353148	1.000000000	0.6177911	0.07370982	-0.003761533	-0.181223
Area	0.61622115	0.617791136	1.0000000	0.17735771	-0.114527566	-0.108972
Anear	0.02055711	0.073709818	0.1773577	1.00000000	-0.119674948	0.044988
Dist	-0.01964207	-0.003761533	-0.1145276	-0.11967495	1.000000000	0.614269
DistSC	-0.18622612	-0.181223511	-0.1089720	0.04498817	0.614269601	1.000000
Elevation	0.70180087	0.746759735	0.7541543	0.56412830	-0.069131928	-0.107342
EM	0.28517661	0.310320605	0.1575627	0.06360495	0.293824516	0.416229

We can see some higher correlations! As to be expected in elevation. This is always important to keep in mind

v. e

```
> gm2 = galamiss[, -2]
> gm2 = gm2[, -7]
> mimgp = amelia(gm2, m=25)
> betasgp=NULL
> sesgp=NULL
> for(i in 1:mimgp$m)
+ {
+   lmod <- lm(NS~Area+Anear+Dist+DistSC+Elevation, mimgp$imputations[[i]])
+   betasgp <- rbind(betasgp, coef(lmod))
+   sesgp <- rbind(sesgp, coef(summary(lmod))[, 2])
+ }
> (cr <- mi.meld(q=betasgp, se=sesgp))
$q.mi
(Intercept)      Area      Anear      Dist      DistSC      Elevation
[1,]      8.010774 -0.02177011 -0.07633469 0.004976406 -0.23399 0.3148571
```

```
$se.mi
```

```
(Intercept)      Area      Anear      Dist      DistSC      Elevation
[1,]      22.45479 0.02413709 0.01935359 1.114134 0.2331302 0.05907911
```

```
> cr$q.mi/cr$se.mi
```

```
(Intercept)      Area      Anear      Dist      DistSC      Elevation
[1,]      0.3567512 -0.9019359 -3.944215 0.004466615 -1.003688 5.329415
```

Looking at our coefficients P values, we can see Elevation and Anear are still the most statistically significant.

(b) 13.3

i. a

```
> pimamiss
```

```
pregnant glucose diastolic triceps insulin  bmi diabetes age test
1          6      148          72      35      NA 33.6      0.627  50
2          1       85          66      29      NA 26.6      0.351  31
3          8      183          64      NA      NA 23.3      0.672  32
4          1       89          66      23      94 28.1      0.167  21
5          0      137          40      35     168 43.1      2.288  33
6          5      116          74      NA      NA 25.6      0.201  30
7          3       78          50      32      88 31.0      0.248  26
8         10      115          NA      NA      NA 35.3      0.134  29

[ reached getOption("max.print") -- omitted 657 rows ]
```

for brevity, a lot of the printouts are removed from this document. You can find the code in the R scripts. I chose not to turn 0s in pregnant and test to NAs as it makes sense for those to be zero. I did let there others turn to NAs, as we can see, insulin levels are tough to get and are often missing. tricep is second most commonly missed, and then diastolic and bmi.

ii. b

```
> summary(pimalm)
```

Call:

```
lm(formula = diastolic ~ ., data = pimamiss)
```

Residuals:

Min	1Q	Median	3Q	Max
-49.420	-6.956	-0.604	7.432	29.268

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	41.004185	4.043536	10.141	< 2e-16 ***
pregnant	0.183487	0.247575	0.741	0.459064
glucose	0.047134	0.025848	1.824	0.069003 .
triceps	-0.005719	0.074506	-0.077	0.938851
insulin	-0.008268	0.006027	-1.372	0.170913
bmi	0.532806	0.112798	4.724	3.26e-06 ***
diabetes	-3.213760	1.722406	-1.866	0.062826 .
age	0.284048	0.081494	3.485	0.000548 ***
test	0.047652	1.508849	0.032	0.974822

Signif. codes: 0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

Residual standard error: 11.38 on 383 degrees of freedom

(376 observations deleted due to missingness)

Multiple R-squared: 0.1882, Adjusted R-squared: 0.1712

F-statistic: 11.1 on 8 and 383 DF, p-value: 3.94e-14

poor fit, and a lot of missing data. We can see bmi and age are important, considering we are missing so much data, it doesn't seem like this model is sufficient. There seems to be some bias to the insulin levels.

iii. c

```
> pimameans = colMeans(pimamiss, na.rm = TRUE)
```

```

> impimamiss = pimamiss
> for(i in c(2:8)) impimamiss[is.na(pimamiss[,i]),i] <- pimameans[i]
> pimaimlm = lm(diastolic~.,impimamiss)
> summary(pimaimlm)

```

Call:

```
lm(formula = diastolic ~ ., data = impimamiss)
```

Residuals:

Min	1Q	Median	3Q	Max
-48.879	-6.599	-0.694	6.369	56.998

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	43.205265	2.620456	16.488	< 2e-16 ***
pregnant	0.157970	0.141327	1.118	0.26402
glucose	0.048453	0.016310	2.971	0.00306 **
triceps	0.006457	0.054022	0.120	0.90489
insulin	-0.007388	0.005139	-1.438	0.15095
bmi	0.476441	0.071163	6.695	4.19e-11 ***
diabetes	-2.127135	1.221251	-1.742	0.08195 .
age	0.285792	0.041421	6.900	1.10e-11 ***
test	-0.868070	1.002583	-0.866	0.38686

Signif. codes: 0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

Residual standard error: 10.92 on 759 degrees of freedom

Multiple R-squared: 0.1938, Adjusted R-squared: 0.1853

F-statistic: 22.8 on 8 and 759 DF, p-value: < 2.2e-16

Mean imputation provides slightly improved results in terms of R squared values.
And glucose has become very significantt.

- iv. d The problem described in the book didnt make sense applied to this problem. I would guess it is asking for us to regress and impute for every column and replace the NA values.

v. e

```
> mimpima = amelia(pima2,m=25)
```

```
> (cr <- mi.meld(q=betaspima,se=sespima))
```

```
$q.mi
```

```
(Intercept)  pregnant    glucose    triceps    insulin      bmi  diabetes
[1,]         40.0757 0.1404388 0.06037759 -0.01842112 -0.008271171 0.5346332 -2.0052
```

```
$se.mi
```

```
(Intercept)  pregnant    glucose    triceps    insulin      bmi  diabetes
[1,]         2.998815 0.1529132 0.02036614 0.06359883 0.006206634 0.0877693 1.270908
```

```
> cr$q.mi/cr$se.mi
```

```
(Intercept)  pregnant    glucose    triceps    insulin      bmi  diabetes    age
[1,]         13.36384 0.918422 2.964606 -0.2896456 -1.332634 6.091347 -1.577835 6.604
```

Looking at the p values, we can see glucose,bmi, and age are quite significantt.

3. Finish R exercises 8.1, 8.2, 8.6, of the textbook. Submit your answers for **ALL** questions.

(a) 8.1

i. a




Figure 7: variance check on pipeline data

clearly there is some fanning here

ii. b

```
> summary(pipwlm)
```

Call:

```
lm(formula = Lab ~ Field, data = pipeline, weights = 1/((Field)^a_1))
```

Weighted Residuals:

Min	1Q	Median	3Q	Max
-1.7450	-0.6789	-0.2672	0.5205	2.8847

Coefficients:

Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-1.49436	0.90707	-1.647	0.102


```
Field          1.20828      0.03488  34.637   <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 0.9795 on 105 degrees of freedom
Multiple R-squared:  0.9195, Adjusted R-squared:  0.9188
F-statistic: 1200 on 1 and 105 DF,  p-value: < 2.2e-16
```

we see some improved R squared values as we diminish the values in order to try and prevent the fanning effect

iii. c




Figure 8: variance check on pipeline data after transform

This is the results of taking the square root on both the response and explanatory variables. It worked quite well.

(b) 8.2

i. a

we can see there is a correlation over time between the residuals/errors




Figure 9: looking at the error correlation of Figure 10: Another look at error correlation
divusa of divusa

ii. b

```
> summary(glusalm)

Generalized least squares fit by maximum likelihood

Model: divorce ~ unemployed + femlab + marriage + birth + military
Data: divusa
AIC      BIC    logLik
179.9523 198.7027 -81.97613

Correlation Structure: AR(1)
Formula: ~year
Parameter estimate(s):
Phi
0.9715486

Coefficients:
Value Std.Error  t-value p-value
(Intercept) -7.059682  5.547193 -1.272658  0.2073
unemployed   0.107643  0.045915  2.344395  0.0219
femlab        0.312085  0.095151  3.279878  0.0016
```

marriage	0.164326	0.022897	7.176766	0.0000
birth	-0.049909	0.022012	-2.267345	0.0264
military	0.017946	0.014271	1.257544	0.2127

Correlation:

(Intr) unmply femlab marrig birth

unemployed -0.420

femlab -0.802 0.240

marriage -0.516 0.607 0.307

birth -0.379 0.041 0.066 -0.094

military -0.036 0.436 -0.311 0.530 0.128

Standardized residuals:

Min	Q1	Med	Q3	Max
-1.4509327	-0.9760939	-0.6164694	1.1375377	2.1593261

Residual standard error: 2.907665

Degrees of freedom: 77 total; 71 residual

> intervals(glusalm,which="var-cov")

Approximate 95% confidence intervals

Correlation structure:

lower	est.	upper
-------	------	-------

Phi 0.6528097 0.9715486 0.9980192

attr("label")

[1] "Correlation structure:"

Residual standard error:

lower	est.	upper
-------	------	-------

0.7974404 2.9076645 10.6020628

we can see that unemployed has become significant, in the previous model, the pvalue was higher.

Further their correlation is significant, we see a positive correlation with a confidence interval that is quite strong

- iii. c Personally, I believe these are correlated over the years mainly due to the warts the data set covers. Baby boomers are all likely to get married around the same time, and thus divorce in similar times as well. Further, War usually causes couples to get married just before leaving for service or after. Thus when they return they will realize they werent meant to be and similarly get divorced at similar times.

(c) 8.6

- i. a




Figure 11: we can see a somewhat linear trend over time that is decreasing.

Not a strong indicator, but something

- ii. b

Generalized least squares fit by REML

Model: taste ~ . - time

Data: c2

AIC BIC logLik

214.94 222.4886 -101.47

Correlation Structure: AR(1)

Formula: ~time

Parameter estimate(s):

Phi

0.2641944

Coefficients:

Value Std.Error t-value p-value

(Intercept) -30.332472 20.273077 -1.496195 0.1466

Acetic 1.436411 4.876581 0.294553 0.7707

H2S 4.058880 1.314283 3.088284 0.0047

Lactic 15.826468 9.235404 1.713674 0.0985

Correlation:

(Intr) Acetic H2S

Acetic -0.899

H2S 0.424 -0.395

Lactic 0.063 -0.416 -0.435

Standardized residuals:

Min Q1 Med Q3 Max

-1.64546468 -0.63861716 -0.06641714 0.52255676 2.41323021

Residual standard error: 10.33276

Degrees of freedom: 30 total; 26 residual

```
> intervals(cgls,which="var-cov")
Approximate 95% confidence intervals
```

```
Correlation structure:
lower      est.      upper
Phi -0.1690265 0.2641944 0.6118599
attr("label")
[1] "Correlation structure:"
```

```
Residual standard error:
lower      est.      upper
7.62646 10.33276 13.99940
```

We can see that the confidence interval include 0, and thus we can not really trust this correlation.

```
iii. > clm2 = lm(taste~.,c2)
> summary(clm2)
```

```
Call:
lm(formula = taste ~ ., data = c2)
```

```
Residuals:
Min      1Q  Median      3Q     Max
-22.3523  -4.9735  -0.5089   4.8531  23.1311
```

```
Coefficients:
Estimate Std. Error t value Pr(>|t|)
(Intercept) -36.6127    17.9845  -2.036  0.05250 .
Acetic        4.1275     4.2556   0.970  0.34139
H2S           3.5387     1.1315   3.127  0.00444 **
Lactic       17.9527     7.7875   2.305  0.02973 *
```

```
time          -0.5459      0.2043  -2.672  0.01306 *
```

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 9.112 on 25 degrees of freedom
```

```
Multiple R-squared:  0.7291, Adjusted R-squared:  0.6858
```

```
F-statistic: 16.83 on 4 and 25 DF,  p-value: 8.205e-07
```

Unlike the GLS, our OLS thinks time is significant! Very funny. However, this is not contradictory, LS and GLS are quite different. This is explained in the next part.

iv. d

in the GLS, we are looking at how correlated the error or noise is over "time", or consecutive entries unlike our ordinary LS. Within the OLS the time value is being included to see how it may provide information on our response. The difference lies within the relations. In OLS it changes the significance and value based on a linear combination within each entry. In residuals, we are only considering the impact of the time variable **AFTER** the coefficients have been established

v. e

if i was told that the entries were not in chronological order, then this would make it purely coincidental that consecutive entries are related, and we should randomize their order to avoid the seemingly correlated entries